

SE350 Final Cheat Sheet

Pre-Midterm Material

OS Definition

OS Roles:

- Referee: resource allocation, isolation, and communication b/w Applications, Users.
- Illusionist: Applications appear to have entire machine; infinite cores, (near) infinite memory, reliable storage and network.
- Glue: System Libraries; Hardware Abstraction

OS Challenges: Reliability & Availability; Security & Privacy; Performance

I/O

Device I/O: Memory Mapped (shared space between DRAM and I/O); Port-Mapped.

Programmed I/O: CPU waits for I/O

Interrupt-Driven I/O: CPU pokes for request, device sends interrupt when done.

DMA: No buffer needed for I/O R/W.

Faster DMA: Buffer Descriptor, Queue of I/O Requests

Other Definitions

PCB: Stores where process is stored in memory, where executable image is, which user runs it, privileges

Hardware Support for Dual Mode: Privileged instr, timer interrupt, memory protection(base & bounds, virtual)

Base & bounds flaws: Fixed heap/stack, no memory sharing, fragmentation, no relative memory addresses

Switching safely: Limited entry (interrupt vector), atomic transfer of control, transparent restartable execution

Reasons to switch to kernel: exception, interrupt, system call, polling

Reasons to switch to User: new process, resume, switch process, user-level upcall

Interrupt Stack: Store registers, Frame pointer, locals, and return address. Kernel stack for each process

Switch Steps: Save SP, execution flags, and inst pointer; Switch onto kernel exception stack; Push those 3 values onto new stack; Optionally save error code; Invoke interrupt handler.

Thread States: init, ready, waiting, running, finished

Preemptive Thread: Can switch anytime.

Cooperative: run without interrupt, explicitly give up CPU; long-running threads can monopolize processor (starvation, non-responsiveness)

Data Stored in TCB: Stack info, saved registers, metadata.

Shared State: Heap, global vars, code

Thread Context Switch: copy current thread registers from processor to TCB.

Copy new thread registers from TCB to processor. Save old threads stack pointer.

Multithreaded Processes:

1. user = kernel thread, kernel does switching.
2. green threads, user level library that does switches. (bad: appears as one process to the kernel, not efficient).
3. scheduler activations, kernel gives processor to user lib, thread lib does switch and scheduling.

Safety: Never enter bad state. **Liveness:** Eventually enters good state.

Shared Objects: can be accessed safely by multiple threads. Has synchronization variables (locks)

Lock: synchronization var that provides mutual exclusion

Condition Variables (CV): a sync object that lets thread efficiently wait for a change in shared state that is protected by lock. (always use in a loop). Memoryless.

Spinlocks: for multiprocessor. Processor waits in loop for lock to become free. (low overhead if held briefly, less than context switch). Deadlock can happen unless all interrupts are disabled.

Semaphore: non negative int val. P(): wait for val \geq 0, then val--. V(): val++, wakes up waiters. Can use like a lock. Better for async IO comm.

Structured Sync: add locks to shared objects. Wait in loop. Use signal/broadcast. Leave shared vars in consistent state.

Uniprocessor Locks: implement by temporarily disable/enable interrupts when acquiring/releasing lock. Move threads to WAITING queue if lock is busy.

Multiprocessor Locks: disable/enable interrupts is not enough. Need atomic read-modify-write instruction, will execute atomically to all other processors (**test_and_set instr**). Use this to implement spin locks.

Readers/Writers Lock: one writer if no readers. Many readers if no writer. `kirito = waitpid(pid)` or just `wait(&kirito)`.

Process: instance of program that is running.

Thread: a *single execution sequence* that represents a *separately schedulable* task

Shell: job control system

Event driven: single thread with event queue.

Multithread: create new thread for each event

Implementations

Synchronization

Uniprocessor Lock

```
Lock::acquire() {
    disableInterrupts();
    if (value == BUSY) {
        waiting.add(myTCB);
        myTCB->state = WAITING;
        next = readyList.remove();
        thread_switch(myTCB, next);
        myTCB->state = RUNNING;
    } else {
        value = BUSY;
    }
    enableInterrupts();
}
```

```
Lock::release() {
    disableInterrupts();
    if (!waiting.Empty()) {
        next = waiting.remove();
        next->state = READY;
        readyList.add(next);
    } else {
        value = FREE;
    }
    enableInterrupts();
}
```

In a Uniprocessor machine, simply need to store TCB of current thread in a global variable

Multiprocessor Lock

```
Lock::acquire() {
    spinLock.acquire();
    if (value == BUSY) {
        waiting.add(myTCB);
        scheduler.suspend(&spinlock);
        // scheduler releases spinlock
    } else {
        value = BUSY;
        spinLock.release();
    }
}
```

```
Lock::release() {
    spinLock.acquire();
    if (!waiting.Empty()) {
        next = waiting.remove();
        scheduler.makeReady(next);
    } else {
        value = FREE;
    }
    spinLock.release();
}
```

```
Sched::suspend(SpinLock *lock) {
    TCB *next;
    disableInterrupts();
    schedSpinLock.acquire();
    spinLock->release();
    myTCB->state = WAITING;
    next = readyList.remove();
    thread_switch(myTCB, next);
    myTCB->state = RUNNING;
    schedSpinLock.release();
    enableInterrupts();
}
```

```
Sched::makeReady(TCB *thread) {
    disableInterrupts();
    schedSpinLock.acquire();
    readyList.add(thread);
    thread->state = READY;
    schedSpinLock.release();
    enableInterrupts();
}
```